

NEURAL NETWORKS

Stefan Jansen

PROGRESSING IN YOUR DATA SCIENCE CAREER

LEARNING OBJECTIVES

- ▶ Understand various types of neural networks
- ▶ Applications of neural networks
- ▶ Apply a neural network model for regression
- ▶ Apply a neural network model for classification

OPENING

ARTIFICIAL NEURAL NETWORKS

OPENING

- ▶ Neural networks were first studied in the 1940s (!) as a model of biological neural networks
- ▶ Many advances since then have improved the ability to train and apply neural networks
- ▶ Good for both classification and regression but difficult to interpret model behaviors
- ▶ Deep learning in the past few years has been highly successful for otherwise difficult problems

OPENING

- ▶ Today we will focus on types of neural networks and their applications, and skip some of the more technical details
- ▶ Specifically we'll skip training neural networks -- there are many methods in various situations and the details can be tedious (but not particularly difficult)
- ▶ Methods include backpropagation, gradient descent, and Hessian-free learning

OPENING

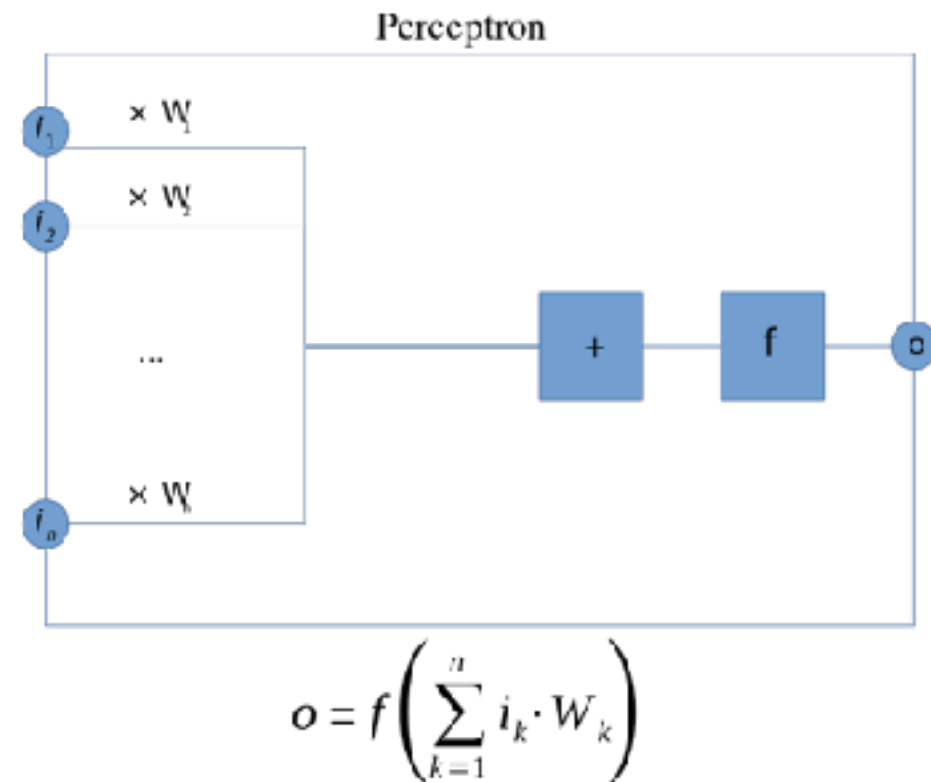

- ▶ Here's a few resources to dive deeper:
 - ▶ [Neural Networks & Deep Learning](#) (by Michael Nielsen)
 - ▶ [The Deep Learning Book](#) (by Goodfellow, Bengio et al)
 - ▶ [Neural Networks for Machine Learning](#) (Geoff Hinton at Coursera)
 - ▶ [Convolutional Neural Networks for Visual Recognition](#)
(Stanford CS 231n)

INTRODUCTION

PERCEPTRON

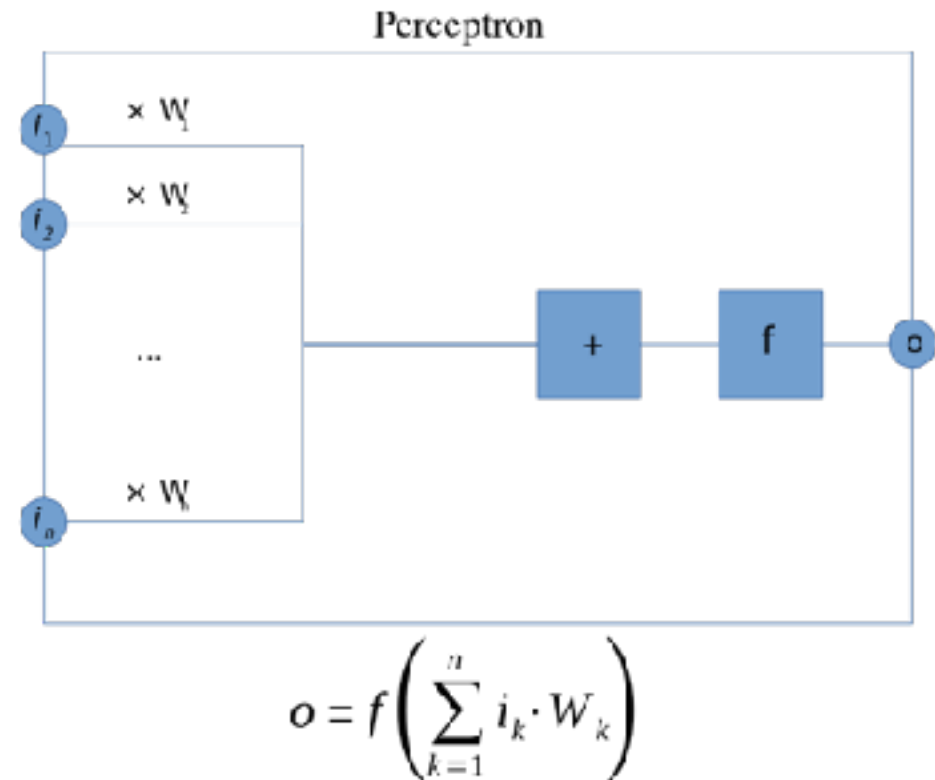

PERCEPTRON

- ▶ Perceptrons are the simplest example of a neural network
- ▶ The idea is to emulate a single neuron



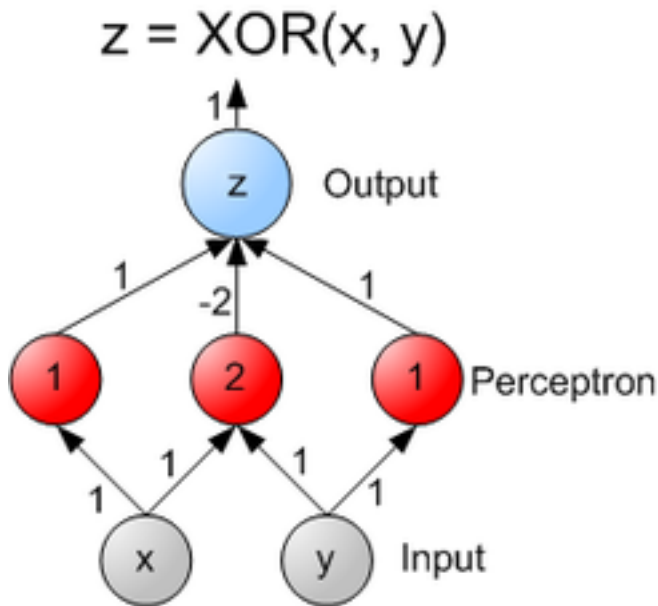
PERCEPTRON

- ▶ Perceptrons are the simplest example of a neural network
- ▶ Given n inputs and an activation or link function f
- ▶ The perceptron computes a linear separating curve




PERCEPTRON

- ▶ Common [activation functions](#) are linear, logistic, tanh, and [softmax](#)
- ▶ We'll see shortly that some are better for classification, some for regression
- ▶ Perceptrons can be combined into multilayer perceptrons or feed-forward network




FEED FORWARD NN

► [Source](#)




FEED FORWARD NN

► [Source](#)




FEED FORWARD NN

► [Source](#)



FEED FORWARD NN

► [Source](#)



FEED FORWARD NN

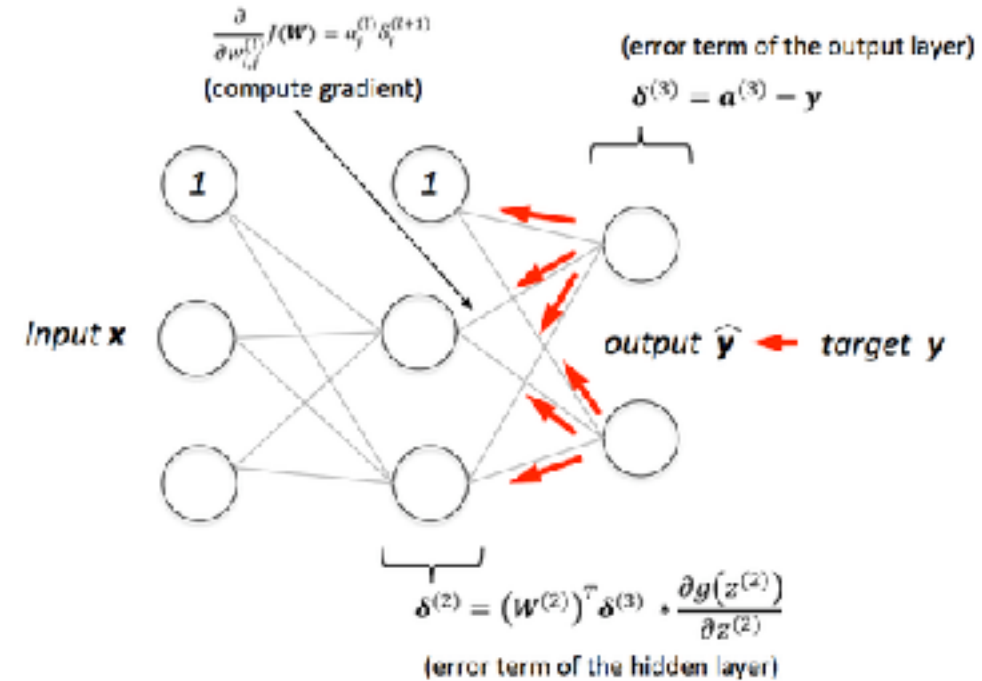

- ▶ Typically we use
 - ▶ Tanh or logistic layers for input
 - ▶ Linear layers for regression output
 - ▶ Logistic or Tanh for binary output
 - ▶ Softmax for n-class output (yields probabilities)

GUIDED PRACTICE

TRAINING

TRAINING

- ▶ Feed forward neural networks can be trained with [backpropagation](#)
- ▶ [Source](#)



TRAINING

- ▶ Key Parameters
 - ▶ Learning Rate (gradient descent for training)
 - ▶ Epochs: number of backpropagation passes (over entire dataset)
 - ▶ Batch size: how many training points used at a time to update weights
- ▶ Model others behaves as usual with
 - ▶ `model.predict`
 - ▶ `model.predict_classes`

TRAINING

► Tips

- If the error jumps around per epoch, decrease the learning rate
- Taking too long to train: use higher learning rate or batch_size
- High error after convergence?
 - More hidden layers / neurons
 - Normalize data or use PCA

INTRODUCTION

UNIVERSAL APPROXIMATION THEORY

UNIVERSAL APPROXIMATION

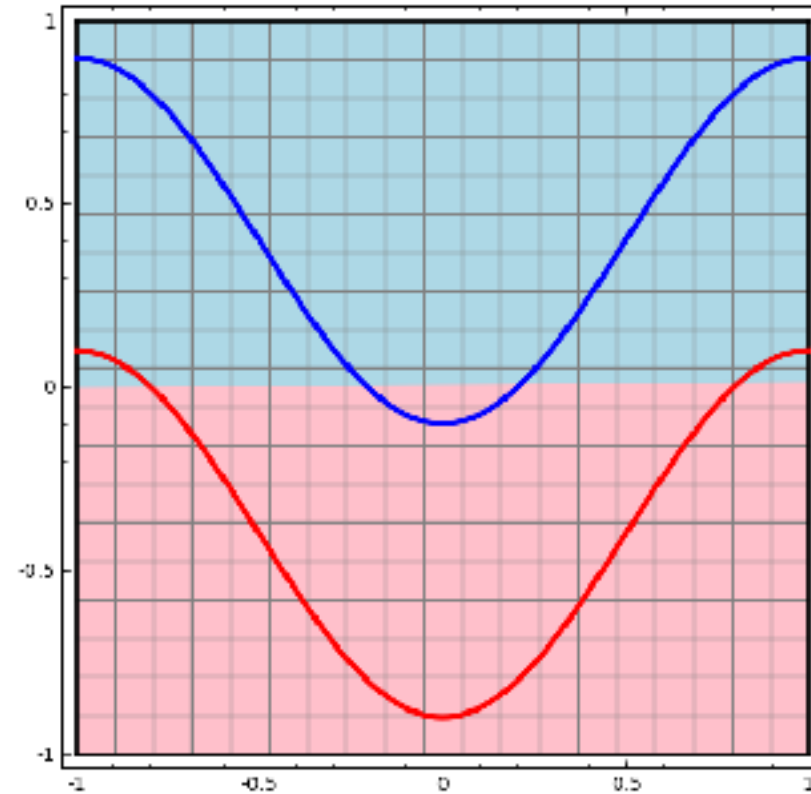
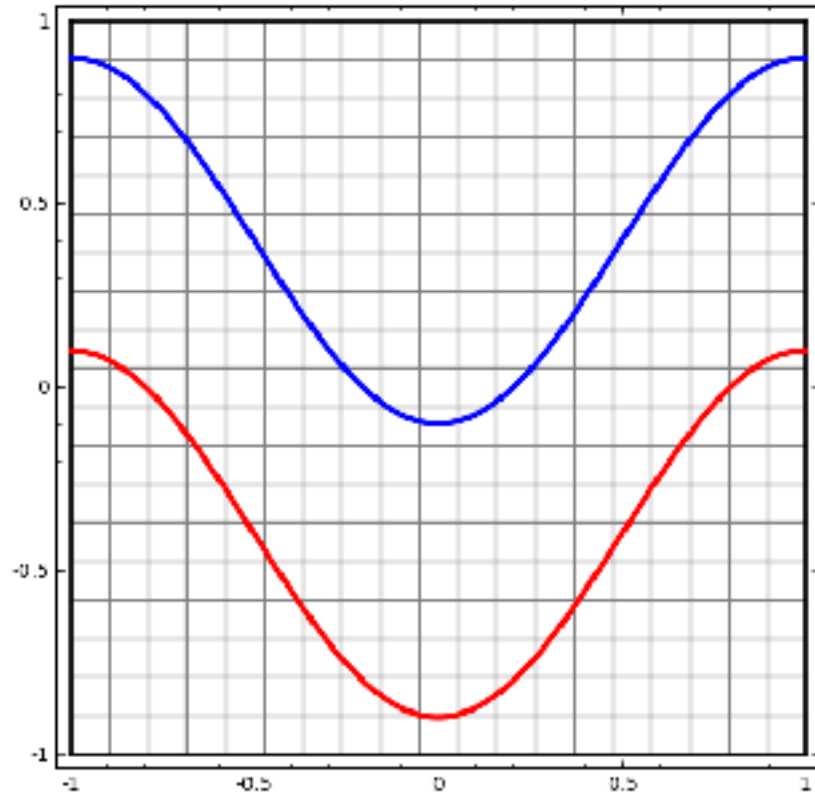
- ▶ One major reason that neural networks are useful is the [Universal Approximation Theorem](#)
- ▶ The result basically says that many real vector-valued functions can be approximated arbitrarily well with *some* feed-forward neural network
- ▶ This is why neural networks are useful for regression -- given enough data and the right network structure they can fit many common data sets

CLASSIFICATION

CLASSIFICATION WITH NEURAL NETWORKS

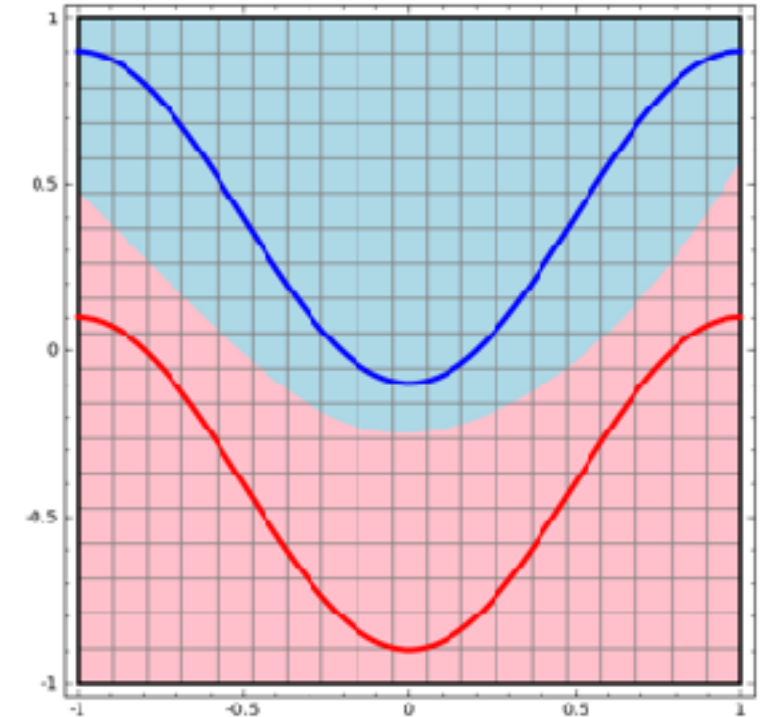


CLASSIFICATION

- ▶ Neural Networks are also extremely useful for classification ([source](#))
- ▶ No hidden layers:



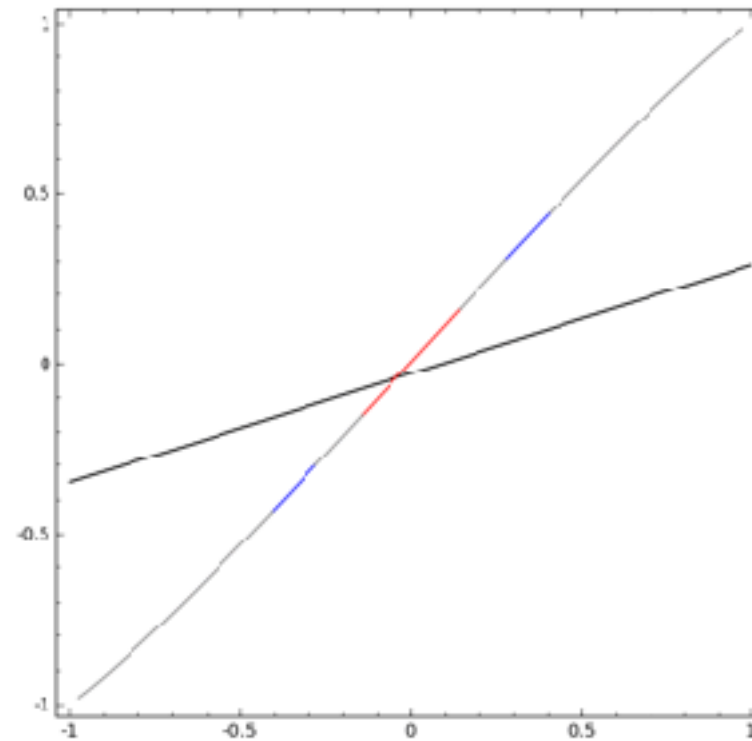
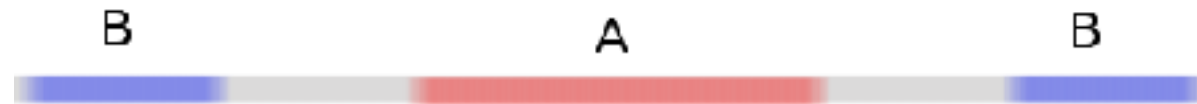
CLASSIFICATION

- ▶ Neural Networks are also extremely useful for classification ([source](#))
- ▶ One hidden layer:



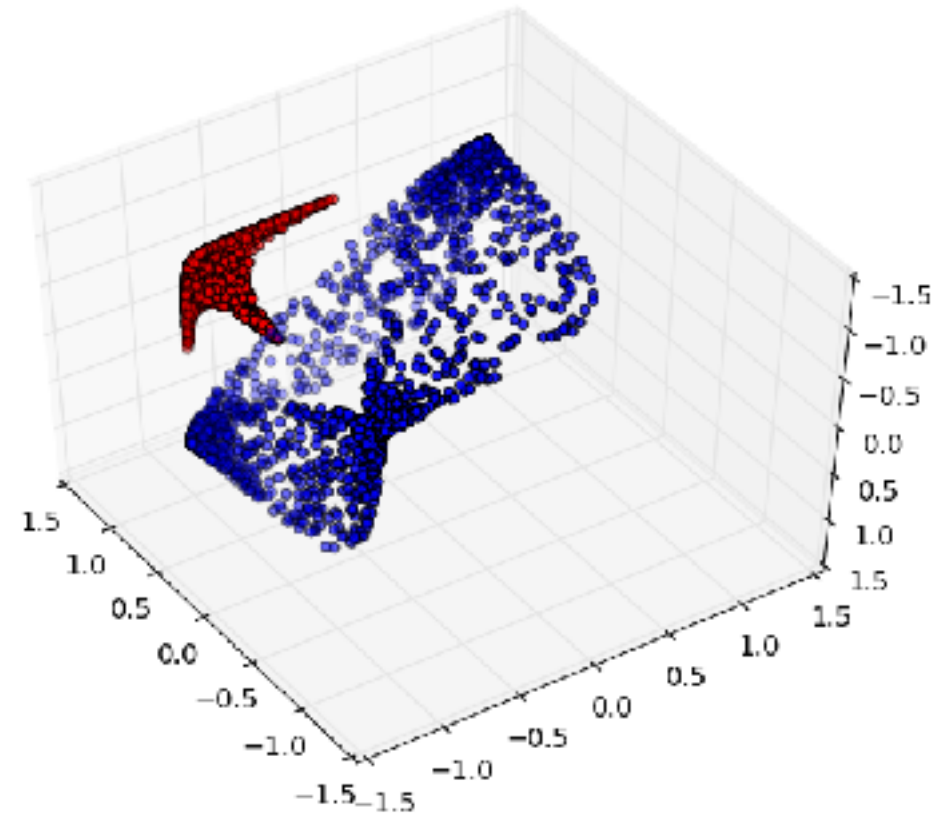

CLASSIFICATION

- Neural Networks are also extremely useful for classification ([source](#))



CLASSIFICATION

- Neural Networks are also extremely useful for classification ([source](#))



CLASSIFICATION

- ▶ The neural network transforms the data topologically (no tears or breaks) and then separates the data with a hyperplane
- ▶ NNs are capable of handling difficult data sets, including:
 - ▶ Image processing: recognizing hand-written characters
 - ▶ Image compression
 - ▶ Financial forecasting
 - ▶ Many others

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



EXERCISE

1. Let's practice using [neural networks for classification](#). For each of the four datasets, experiment with the number of layers and neurons to find the best model
2. Also take a look at this [visualization](#)

DELIVERABLE

Answers to the above questions

GUIDED PRACTICE

NEURAL NETWORKS IN PYTHON

NN IN PYTHON

- There are many NN libraries for python and other languages
- Python
 - [Theano](#)
 - [Keras](#)
 - [Lasagne](#)
 - [TensorFlow](#)
 - Scikit Learn support for NN coming in 0.18
- Lua
 - [Torch](#)
- Some of these libraries utilize GPUs for (much) faster training

NN IN PYTHON

- ▶ Let's look at some examples in Keras
 - ▶ Regression
 - ▶ Classification

GUIDED PRACTICE

DESIGNING NEURAL NETWORKS

NN IN PYTHON


- ▶ Network design is a hard problem
 - ▶ Experience helps
 - ▶ Evolutionary algorithms are [useful](#) for [design](#)
 - ▶ Nice (free) book [available](#)

RECURRENT NN

RECURRENT NEURAL NETWORKS

RECURRENT NEURAL NETWORKS

- ▶ Recurrent Neural Networks contain loops ([source](#))



RECURRENT NEURAL NETWORKS

- ▶ Recurrent Neural Networks contain loops
- ▶ This implements feedback and gives neural networks “memory” or context
- ▶ Particularly good for predicting sequences, translating text, recognizing objects in images, speech translation
- ▶ Commonly referred to as **deep learning**, involving both feature extraction and modeling
- ▶ [Nice intro here](#)

RECURRENT NEURAL NETWORKS


► [RNN font analysis](#)



A	B	C	D	E	F	G	H
I	J	K	L	M	N	O	P
Q	R	S	T	U	V	W	X
Y	Z	a	b	c	d	e	f
g	h	i	j	k	l	m	n
o	p	q	r	s	t	u	v
w	x	y	z	0	1	2	3
4	5	6	7	8	9		

RECURRENT NEURAL NETWORKS

► Automatic Colorization with CNN



RECURRENT NEURAL NETWORKS

- ▶ [RNN font analysis](#)
- ▶ [Automatic Colorization](#) with CNN
- ▶ Automatic translation
- ▶ [Deep Learning Applications](#)

CONCLUSION

TOPIC REVIEW

CONCLUSION: Neural Networks

Pros:

- Flexible
- Good for a variety of tasks
- Good for many types of data

Cons:

- Can require a lot of data
- Training may be slow
- Many parameters to tune
- Many layer types and activations
- Black Box model

CONCLUSION

- ▶ Many [more examples](#) for Keras available
- ▶ Recommended articles: [Convolutional NN](#),
- ▶ Advanced machine learning methods you should explore include Bayesian methods and deep learning

LESSON

Q & A

LESSON

EXIT TICKET

**DON'T FORGET TO FILL OUT YOUR EXIT
TICKET**